

Pstat 174 Boston Crime Project

Riley Mault

February 18, 2020

1. Abstract

Crime incident reports are provided by Boston Police Department (BPD) to document the initial details surrounding an incident to which BPD officers respond. This is a dataset containing records capturing the type of incident as well as when and where it occurred. For this time series project, I will be focusing exclusively on when crimes occurred. Records begin in June 14, 2015 and continue to September 3, 2018.

I will be addressing questions such as: Does the frequency of crimes have any pattern? If the frequency of crimes does have a pattern, why is that? Is it possible to forecast the daily frequency of crimes?

To address these questions, I plotted the data ordered by time, analyzed auto-correlation plots, and developed seasonal ARIMA models to predict/forecast future values.

2.

Introduction

My goal of this daily Boston crime report dataset is to find and observe frequencies at which crimes occur. In observing those frequencies, I will build an ARIMA model to forecast the frequencies of future crime reports.

Importing and Cleaning Data

```
# import dataset
crime_df = read.csv('crime.csv')
attach(crime_df)
```


```
# sorting data by date of crime
crime_df = crime_df %>% separate(OCCURRED_ON_DATE, c("Date", "Time"), sep = " ") %>%
mutate(Date = ymd(Date))
crime_df$Date = as.Date(crime_df$Date)
crime_df = crime_df[order(crime_df$Date),]
```

Distribution of the Counts of Crimes by Date

```
ts_data_full = ts(table(crime_df$Date))

# split up data into train/test split
ts_data = ts(ts_data_full[c(1:1067)])
ts_data.test = ts(ts_data_full[c(1068:1077)])

# histogram of our crime data
hist(ts_data, 50, col='light blue', main='Crimes Count Distribution')
```



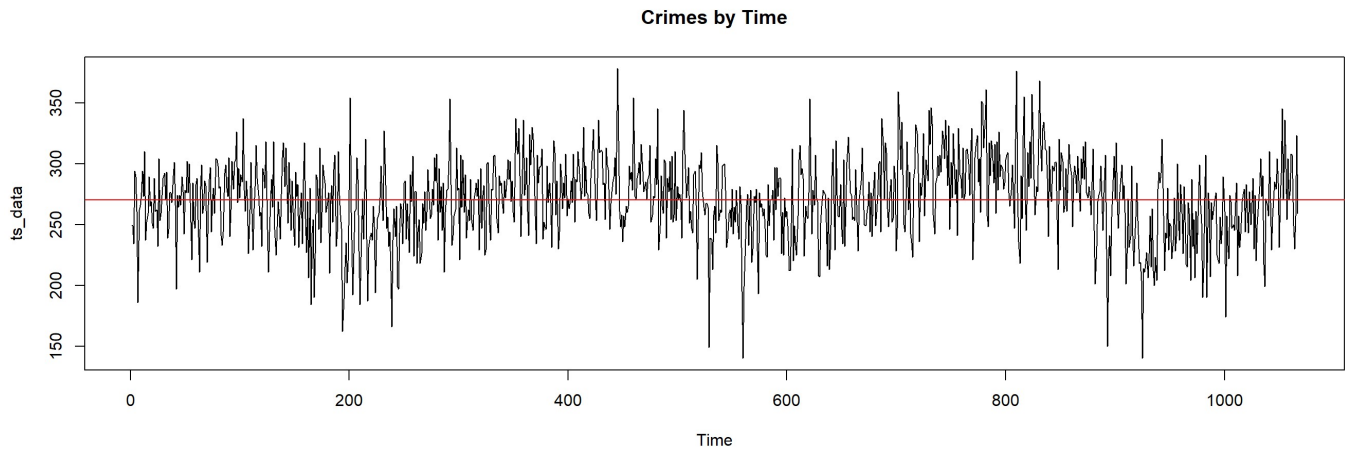
```
paste(c('Skewness is', skewness(ts_data)), collapse = ' ')
```

```
## [1] "Skewness is -0.22456247991086"
```

I separated the data into 50 bins and plotted a histogram. Most days, about 250 - 300 crimes occur. We can see that the data seems pretty normal with a very slight left skew. Therefore, a transformation isn't necessary.

Distribution of Crimes by Time

```
# Plotting timeseries plot
ts.plot(ts_data, main = "Crimes by Time")
abline(h=mean(ts_data), col='red')
```



```
# Transformation of data but not needed/performed

# bcTransform<-boxcox(ts_data ~ as.numeric(1:length(ts_data)))
# bcTransform$x[which(bcTransform$y== max(bcTransform$y))]
# lambda = bcTransform$x[which(bcTransform$y== max(bcTransform$y))]
# lambda = bcTransform$x[which(bcTransform$y== max(bcTransform$y))]
# ts.bc = (1/lambda)*(ts_data^lambda-1)
# hist(ts.bc, 50, col='light blue', main='Crimes Count Distribution')
# shapiro.test(ts.bc)
```


From the chart above, there is noticeable seasonality that we will address in order to make our data stationary. We can see a repeating pattern (almost that of a sin wave) every 365 days.

ACF and PACF


```
par(mfrow=c(1,2))

# ACF and PACF
acf(table(crime_df$Date), 40, main='Autocorrelation Lag=40')
pacf(table(crime_df$Date), 40, main='Partial Autocorrelation Lag=40')
```

Autocorrelation Lag=40




Partial Autocorrelation Lag=40



In both the ACF and PACF plots, there are lag spikes every 7 lags. To approach this, I will difference the data by 7 lag units. The partial correlation shows many significant lags. We can conclude that crimes are correlated with yesterday and the same day in each previous week.

```
# Decompose ts to observe trend/seasonality
y = ts(as.ts(ts_data), frequency=7)
plot(decompose(y))
```

Decomposition of additive time series

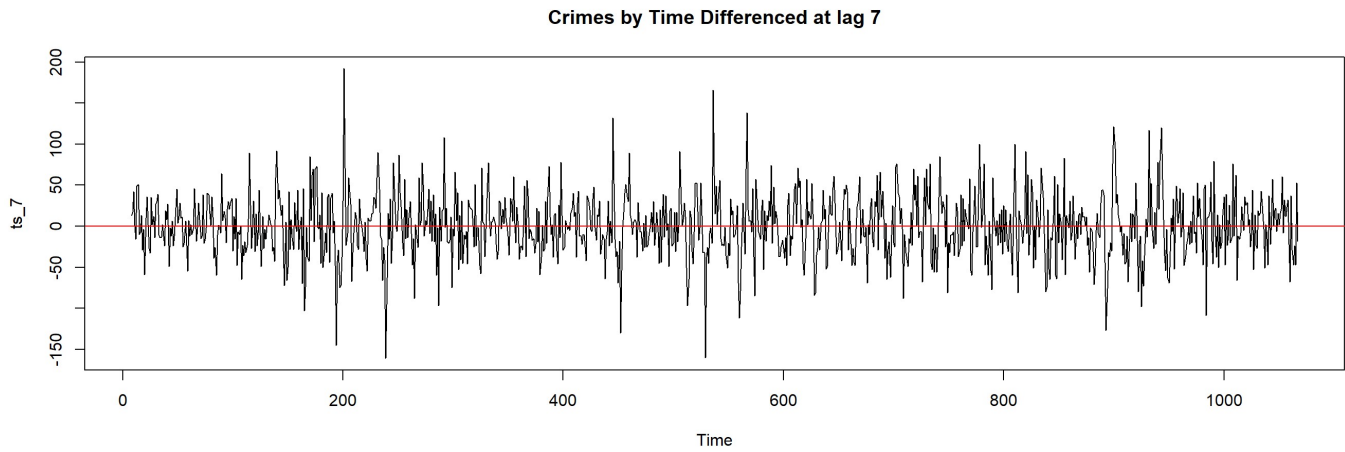


In decomposing the crime data, we notice a small amount of trend and noticeable seasonality. Again, this gives us more reason to difference the data to make it stationary.

Differencing Data

```
# Differencing the model by 7 to remove seasonality/trend
ts_7 = diff(ts_data, lag=7)

plot.ts(ts_7, main="Crimes by Time Differenced at lag 7")
abline(h=mean(ts_7), col='red')
```



Dickey-Fuller Test for Stationary

```
adf.test(ts_7)
```

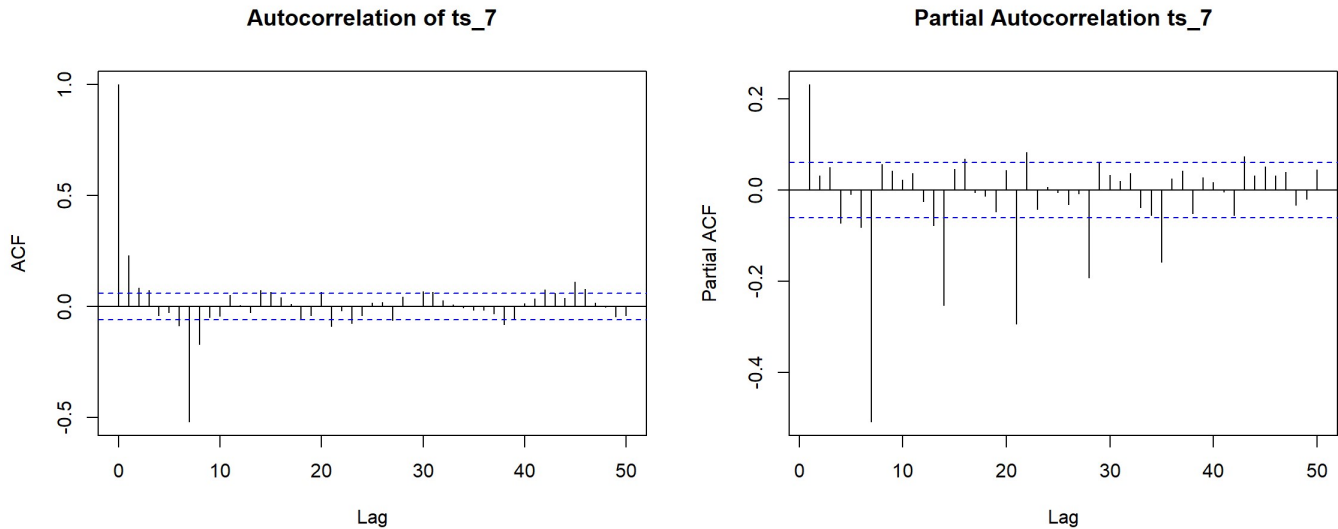
```
## Warning in adf.test(ts_7): p-value smaller than printed p-value
```

```
##  
## Augmented Dickey-Fuller Test  
##  
## data: ts_7  
## Dickey-Fuller = -11.633, Lag order = 10, p-value = 0.01  
## alternative hypothesis: stationary
```

After differencing the data by 7, the time series graph looks much more stationary. The p-value of the Dickey-Fuller Test (smaller than .01) being less than .05 confirms the alternative hypothesis that the data is stationary.

Differenced ACF and PACF

```
par(mfrow=c(1,2))  
  
acf(ts_7, 50, main='Autocorrelation of ts_7')  
pacf(ts_7,50, main='Partial Autocorrelation ts_7')
```



After differencing our data, the ACF and PACF plots look much cleaner and easier to analyze. In the ACF plot the first few lags show significance. From this, we can possibly assign a q value of 2 or 3 for the non-seasonal component. There is another significant spike at 7 due to the seasonal component; so we assign $Q=1$.

In the PACF plot, there is one significant lag at lag 1. However, we might be better off having a complete MA model and having no AR components to the non-seasonal aspect. There are also continuous decreasing lags every 7 lags. From this, can attribute a seasonal AR component with order of at least 3 or 4.

ARIMA Models

```
# Model A
fit.A = arima(ts_data, order=c(0,1,3), seasonal = list(order = c(4,1,1), period = 7),
method="CSS")
fit.A
```

```
##
## Call:
## arima(x = ts_data, order = c(0, 1, 3), seasonal = list(order = c(4, 1, 1), period =
7),
##     method = "CSS")
##
## Coefficients:
##          ma1          ma2          ma3          sar1          sar2          sar3          sar4
##      -0.7326  -0.1575  -0.0236  -0.065  -0.0066  -0.1432  -0.0179
## s.e.   0.0310   0.0403   0.0307   0.032   0.0274   0.0308   0.0301
##          sma1
##      -0.9644
## s.e.   0.0090
##
## sigma^2 estimated as 734.4:  part log likelihood = -4996.88
```

```
AICc(arima(ts_data, order=c(0,1,3), seasonal = list(order = c(4,1,1), period = 7), me
thod="ML"))
```

```
## [1] 9997.219
```

```
# Model B
fit.B = arima(ts_data, order=c(0,1,2), seasonal = list(order = c(4,1,1), period = 7),
method="CSS")
fit.B
```

```
##
## Call:
## arima(x = ts_data, order = c(0, 1, 2), seasonal = list(order = c(4, 1, 1), period =
7),
##      method = "CSS")
##
## Coefficients:
##          ma1          ma2          sar1          sar2          sar3          sar4          sma1
##      -0.7347  -0.1771  -0.0677  -0.0090  -0.1462  -0.0183  -0.9642
## s.e.   0.0303   0.0303   0.0319   0.0274   0.0305   0.0301   0.0090
##
## sigma^2 estimated as 734.8:  part log likelihood = -4997.18
```

```
AICc(arima(ts_data, order=c(0,1,2), seasonal = list(order = c(4,1,1), period = 7), met
hod="ML"))
```

```
## [1] 9995.625
```

After testing out many similar models to the ps and qs suggested from the ACF/PACF plots, these two models (Model A and Model B) resulted in the lowest AICc values and best fitting models. Model A has MA order 3 coefficients while Model B has MA order 2. Model B has a slightly lower AICc value by 2. To decide which model is better, we will run diagnostics on both.


Diagnostic Checking for Model A

$$\Delta_7$$

$$U_t$$


$$= (1 - 0.7326B - 0.1575B^2 - 0.0236B^3)(1 - 0.9644B^7)Z_t + (0.065B^7 - 0.0066B^{14} - 0.1432B^{21} - 0.0179B^{28})X_t$$

```
# Checking Invertibility of model A
autoplot(fit.A)
```




```
# residual hist
res = residuals(fit.A)
hist(res,density=20,breaks=20, col="blue", xlab="", prob=TRUE)
m = mean(res)
std = sqrt(var(res))
curve( dnorm(x,m,std), add=TRUE )
```

Histogram of res





```
plot.ts(res)
fitt = lm(res ~ as.numeric(1:length(res))); abline(fitt, col="red")
abline(h=mean(res), col="blue")
```




```
# Normal Q-Q to check for normality of res
qqnorm(res, main= "Normal Q-Q Plot for Model A")
qqline(res, col="blue")
```


Normal Q-Q Plot for Model A



```
acf(res, lag.max=40)
```



```
pacf(res, lag.max=40)
```



```
# Tests to check fit  
df = 3  
shapiro.test(res)
```

```
##
##  Shapiro-Wilk normality test
##
## data:  res
## W = 0.98705, p-value = 4.116e-08
```

```
Box.test(res, lag = 7, type = c("Box-Pierce"), fitdf= df)
```

```
##
##  Box-Pierce test
##
## data:  res
## X-squared = 2.0956, df = 4, p-value = 0.7182
```

```
Box.test(res, lag = 7, type = c("Ljung-Box"), fitdf= df)
```

```
##
##  Box-Ljung test
##
## data:  res
## X-squared = 2.1069, df = 4, p-value = 0.7161
```

```
Box.test(res^2, lag = 7, type = c("Ljung-Box"), fitdf= 0)
```

```
##
##  Box-Ljung test
##
## data:  res^2
## X-squared = 12.243, df = 7, p-value = 0.09285
```


Diagnostic Checking for Model B

$$\Delta_7$$

$$U_t$$


$$= (1 - 0.7347B - 0.1771B^2)(1 - 0.9642B^7)Z_t + (-0.0677B^7 - 0.0090B^{14} - 0.1462B^{21} - 0.0183B^{28})X_t$$

```
#Checking Invertibility of model B
autoplot(fit.B)
```




```
# residual hist
res = residuals(fit.B)
hist(res,density=20,breaks=20, col="blue", xlab="", prob=TRUE)
m = mean(res)
std = sqrt(var(res))
curve( dnorm(x,m,std), add=TRUE )
```

Histogram of res




```
plot.ts(res)
fitt = lm(res ~ as.numeric(1:length(res))); abline(fitt, col="red")
abline(h=mean(res), col="blue")
```




```
# Normal Q-Q to check for normality of res
qqnorm(res, main= "Normal Q-Q Plot for Model B")
qqline(res, col="blue")
```


Normal Q-Q Plot for Model B



```
acf(res, lag.max=40)
```



```
pacf(res, lag.max=40)
```



```
# Tests to check fit  
df = 2  
shapiro.test(res)
```

```
##
##  Shapiro-Wilk normality test
##
## data:  res
## W = 0.98707, p-value = 4.212e-08
```

```
Box.test(res, lag = 7, type = c("Box-Pierce"), fitdf= df)
```

```
##
##  Box-Pierce test
##
## data:  res
## X-squared = 2.4736, df = 5, p-value = 0.7805
```

```
Box.test(res, lag = 7, type = c("Ljung-Box"), fitdf= df)
```

```
##
##  Box-Ljung test
##
## data:  res
## X-squared = 2.4861, df = 5, p-value = 0.7786
```

```
Box.test(res^2, lag = 7, type = c("Ljung-Box"), fitdf= 0)
```

```
##
##  Box-Ljung test
##
## data:  res^2
## X-squared = 12.421, df = 7, p-value = 0.08754
```

Our diagnostics for both models both pretty solid. Both models pass all the necessary tests of fit (since the p-values are all greater than 0.05), except the Shapiro-Wilk normality test. Both models fail the normality test. However, when looking at the histogram and normal Q-Q plot of the residuals, we see that they all look pretty good and normal. Even after attempting to make a model using box-cox transformed data, it still fails the residual normality test. So even though we have this test error issue, the histogram and normal Q-Q plots give us solid evidence that our model is normally distributed. All the roots for both models are within the unit circle. And lastly, the ACF/PACF plots of the residuals look great. Only 2 at max protrude out of the 95% interval boundary.

Forecasting Model A


```
# Predicting future values
forecast(fit.A)
```

##	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
## 1068	274.4687	239.7380	309.1994	221.3526	327.5848
## 1069	291.1030	255.1523	327.0537	236.1211	346.0849
## 1070	272.1628	236.0101	308.3156	216.8720	327.4537
## 1071	245.4002	209.1233	281.6770	189.9195	300.8809
## 1072	272.3563	235.9558	308.7569	216.6865	328.0262
## 1073	277.5209	240.9971	314.0448	221.6625	333.3793
## 1074	276.4096	239.7629	313.0563	220.3633	332.4559
## 1075	282.8314	246.1315	319.5313	226.7037	338.9590
## 1076	289.0594	252.2584	325.8603	232.7772	345.3415
## 1077	269.3490	232.4351	306.2629	212.8941	325.8039
## 1078	241.6142	204.5858	278.6426	184.9841	298.2443
## 1079	267.1865	230.0439	304.3291	210.3818	323.9912
## 1080	283.2895	246.0330	320.5459	226.3107	340.2682
## 1081	284.4475	247.0776	321.8174	227.2952	341.5998


```

pred.A = predict(fit.A, n.ahead= 110)
U.A= pred.A$pred+ 1.96*pred.A$se
L.A= pred.A$pred-1.96*pred.A$se
ts.plot(ts_data, xlim=c(950,length(ts_data)+110))
lines(U.A, col="blue", lty=2)
lines(L.A, col="blue", lty=2)
lines(1068:1177, pred.A$pred, col="red")

```




```
# Adding original data to predictions
ts.plot(ts_data_full, xlim= c(950,length(ts_data)+110), col="black")
lines(U.A, col="blue", lty="dashed")
lines(L.A, col="blue", lty="dashed")
lines((length(ts_data)+1):(length(ts_data)+110), pred.A$pred, col="red")
lines((length(ts_data)+1):(length(ts_data)+110), pred.A$pred, col="red")
```



Forecasting Model B


```
# Predicting future values
forecast(fit.B)
```

##	Point	Forecast	Lo 80	Hi 80	Lo 95	Hi 95
## 1068	273.7099	238.9695	308.4503	220.5791	326.8407	
## 1069	292.0265	256.0847	327.9682	237.0583	346.9946	
## 1070	272.3912	236.3193	308.4630	217.2240	327.5583	
## 1071	245.6800	209.4785	281.8815	190.3146	301.0454	
## 1072	272.5438	236.2131	308.8745	216.9808	328.1068	
## 1073	277.6559	241.1965	314.1153	221.8960	333.4157	
## 1074	276.5554	239.9678	313.1431	220.5994	332.5114	
## 1075	283.0766	246.4369	319.7163	227.0410	339.1122	
## 1076	289.1334	252.3894	325.8774	232.9382	345.3286	
## 1077	269.5850	232.7216	306.4483	213.2073	325.9626	
## 1078	241.8360	204.8537	278.8182	185.2765	298.3955	
## 1079	267.3441	230.2434	304.4449	210.6034	324.0849	
## 1080	283.4828	246.2638	320.7017	226.5613	340.4042	
## 1081	284.7771	247.4404	322.1139	227.6756	341.8787	

```

pred.B = predict(fit.B, n.ahead= 110)
U.B= pred.B$pred+ 1.96*pred.B$se
L.B= pred.B$pred-1.96*pred.B$se
ts.plot(ts_data, xlim=c(950,length(ts_data)+110))
lines(U.B, col="blue", lty=2)
lines(L.B, col="blue", lty=2)
lines(1068:1177, pred.B$pred, col="red")


```



```

# Adding original data to predictions
ts.plot(ts_data_full, xlim= c(950,length(ts_data)+110), col="black")
lines(U.B, col="blue", lty="dashed")
lines(L.B, col="blue", lty="dashed")
lines((length(ts_data)+1):(length(ts_data)+110), pred.B$pred, col="red")
lines((length(ts_data)+1):(length(ts_data)+110), pred.B$pred, col="red")

```



After examining the forecast predictions for both models, we can see that they are extremely similar. They both show somewhat accurate predictions and do a good job articulating the weekly seasonality. As it attempts to predict more, it very slightly begins to over/under predict possibly due to the changing yearly seasonality (Not uncommon to see in models such as these). Because it is very difficult to notice differences, we will choose Model B being better based on its lower AICc value.

3. Conclusion

After analyzing the time series and acf plots of the crime data, it is apparent that there is a pattern in the frequency of crimes. It seems that there is a yearly pattern as well as a weekly pattern. Friday appears to have the highest number of crimes per week, while Sunday the least. From the data, it is not quite clear why most crimes occur on Friday. Maybe we can attribute it to Friday night being the most popular night of the week to leave the house; And Sunday the least likely night of the week to leave the house. However, there are many parameters that can attribute to this, and we need more outside data to identify the reason.

Due to the strong seasonality of the data, it was necessary to use a SARIMA model in order to forecast future data. From understanding the weekly seasonal pattern, I was able to make a couple SARIMA models that accurately forecast the daily frequency of crimes. Although the residuals of both models appeared normal in the diagnostic plots, it is important to note that they did not pass the Shapiro-Wilk normality test.

4. References

Dataset <https://www.kaggle.com/AnalyzeBoston/crimes-in-boston> (<https://www.kaggle.com/AnalyzeBoston/crimes-in-boston>)

Skewness <https://www.r-bloggers.com/measures-of-skewness-and-kurtosis/> (<https://www.r-bloggers.com/measures-of-skewness-and-kurtosis/>)

Decompose <http://r-statistics.co/Time-Series-Analysis-With-R.html> (<http://r-statistics.co/Time-Series-Analysis-With-R.html>)

With-R.html)

Adf Stationary Testing <https://nwfsc-timeseries.github.io/atsa-labs/sec-boxjenkins-aug-dickey-fuller.html>
(<https://nwfsc-timeseries.github.io/atsa-labs/sec-boxjenkins-aug-dickey-fuller.html>)

Forecast <https://robjhyndman.com/hyndsight/forecast7-ggplot2/> (<https://robjhyndman.com/hyndsight/forecast7-ggplot2/>)

5. Appendix

```

# import dataset
crime_df = read.csv('crime.csv')
attach(crime_df)

# sorting data by date of crime
crime_df = crime_df %>% separate(OCCURRED_ON_DATE, c("Date", "Time"), sep = " ") %>%
mutate(Date = ymd(Date))
crime_df$Date = as.Date(crime_df$Date)
crime_df = crime_df[order(crime_df$Date),]

ts_data_full = ts(table(crime_df$Date))

# split up data into train/test split
ts_data = ts(ts_data_full[c(1:1067)])
ts_data.test = ts(ts_data_full[c(1068:1077)])

# histogram of our crime data
hist(ts_data, 50, col='light blue', main='Crimes Count Distribution')
paste(c('Skewness is', skewness(ts_data)), collapse = ' ')

# Plotting timeseries plot
ts.plot(ts_data, main = "Crimes by Time")
abline(h=mean(ts_data), col='red')

# Transformation of data but not needed/performed

# bcTransform<-boxcox(ts_data ~ as.numeric(1:length(ts_data)))
# bcTransform$x[which(bcTransform$y== max(bcTransform$y))]
# lambda = bcTransform$x[which(bcTransform$y== max(bcTransform$y))]
# lambda = bcTransform$x[which(bcTransform$y== max(bcTransform$y))]
# ts.bc = (1/lambda)*(ts_data^lambda-1)
# hist(ts.bc, 50, col='light blue', main='Crimes Count Distribution')
# shapiro.test(ts.bc)

# ACF and PACF
acf(table(crime_df$Date), 40, main='Autocorrelation Lag=40')
pacf(table(crime_df$Date), 40, main='Partial Autocorrelation Lag=40')

# Decomposing to examine seasonality/trend
y = ts(as.ts(ts_data), frequency=7)
plot(decompose(y))

# Differencing the model by 7 to remove seasonality/trend
ts_7 = diff(ts_data, lag=7)

plot.ts(ts_7, main="Crimes by Time Differenced at lag 7")
abline(h=mean(ts_7), col='red')

```

```

adf.test(ts_7)

acf(ts_7, 50, main='Autocorrelation of ts_7')
pacf(ts_7, 50, main='Partial Autocorrelation ts_7')

# Model A
fit.A = arima(ts_data, order=c(0,1,3), seasonal = list(order = c(4,1,1), period = 7),
method="CSS")
fit.A
AICc(arima(ts_data, order=c(0,1,3), seasonal = list(order = c(4,1,1), period = 7), me
thod="ML"))

# Model B
fit.B = arima(ts_data, order=c(0,1,2), seasonal = list(order = c(4,1,1), period = 7),
method="CSS")
fit.B
AICc(arima(ts_data, order=c(0,1,2), seasonal = list(order = c(4,1,1), period = 7), met
hod="ML"))

# Checking Invertibility of model A
autoplot(fit.A)

# residual hist
res = residuals(fit.A)
hist(res,density=20,breaks=20, col="blue", xlab="", prob=TRUE)
m = mean(res)
std = sqrt(var(res))
curve( dnorm(x,m,std), add=TRUE )
plot.ts(res)
fitt = lm(res ~ as.numeric(1:length(res))); abline(fitt, col="red")
abline(h=mean(res), col="blue")

# Normal Q-Q to check for normality of res
qqnorm(res,main= "Normal Q-Q Plot for Model A")
qqline(res,col="blue")

acf(res, lag.max=40)
pacf(res, lag.max=40)

# Tests to check fit
df = 3
shapiro.test(res)
Box.test(res, lag = 7, type = c("Box-Pierce"), fitdf= df)
Box.test(res, lag = 7, type = c("Ljung-Box"), fitdf= df)
Box.test(res^2, lag = 7, type = c("Ljung-Box"), fitdf= 0)

#Checking Invertibility of model B
autoplot(fit.B)

# residual hist
res = residuals(fit.B)

```

```

hist(res,density=20,breaks=20, col="blue", xlab="", prob=TRUE)
m = mean(res)
std = sqrt(var(res))
curve( dnorm(x,m,std), add=TRUE )
plot.ts(res)
fitt = lm(res ~ as.numeric(1:length(res))); abline(fitt, col="red")
abline(h=mean(res), col="blue")

# Normal Q-Q to check for normality of res
qqnorm(res,main= "Normal Q-Q Plot for Model B")
qqline(res,col="blue")

acf(res, lag.max=40)
pacf(res, lag.max=40)

# Tests to check fit
df = 2
shapiro.test(res)
Box.test(res, lag = 7, type = c("Box-Pierce"), fitdf= df)
Box.test(res, lag = 7, type = c("Ljung-Box"), fitdf= df)
Box.test(res^2, lag = 7, type = c("Ljung-Box"), fitdf= 0)

# Predicting future values
forecast(fit.A)
pred.A = predict(fit.A, n.ahead= 110)
U.A= pred.A$pred+ 1.96*pred.A$se
L.A= pred.A$pred-1.96*pred.A$se
ts.plot(ts_data, xlim=c(950,length(ts_data)+110))
lines(U.A, col="blue", lty=2)
lines(L.A, col="blue", lty=2)
lines(1068:1177, pred.A$pred, col="red")

# Adding original data to predictions
ts.plot(ts_data_full, xlim= c(950,length(ts_data)+110), col="black")
lines(U.A, col="blue", lty="dashed")
lines(L.A, col="blue", lty="dashed")
lines((length(ts_data)+1):(length(ts_data)+110), pred.A$pred, col="red")
lines((length(ts_data)+1):(length(ts_data)+110), pred.A$pred, col="red")

# Predicting future values
forecast(fit.B)
pred.B = predict(fit.B, n.ahead= 110)
U.B= pred.B$pred+ 1.96*pred.B$se
L.B= pred.B$pred-1.96*pred.B$se
ts.plot(ts_data, xlim=c(950,length(ts_data)+110))
lines(U.B, col="blue", lty=2)
lines(L.B, col="blue", lty=2)
lines(1068:1177, pred.B$pred, col="red")

# Adding original data to predictions
ts.plot(ts_data_full, xlim= c(950,length(ts_data)+110), col="black")

```

```
lines(U.B, col="blue", lty="dashed")
lines(L.B, col="blue", lty="dashed")
lines((length(ts_data)+1):(length(ts_data)+110), pred.B$pred, col="red")
lines((length(ts_data)+1):(length(ts_data)+110), pred.B$pred, col="red")
```